

Programming Assignment 4

Due May 12, 2022 at 11:59pm

For this assignment and PA5, you will implement **the intermediate code generation phase** of your *Cool* compiler. In PA4, you will implement only the language features of *Cool* shown in the figure below.

1 COOL Language Subset for PA4

```

program ::= class;
class ::= class Main { feature; }
feature ::= main() : Int { expr }
expr ::= ID <- expr
        | if expr then expr else expr fi
        | while expr loop expr pool
        | { [expr;]+ }
        | let ID : TYPE [ <- expr ] [ID : TYPE [ <- expr ] ]* in expr
        | expr + expr
        | expr - expr
        | expr * expr
        | expr / expr
        | ~ expr
        | expr < expr
        | expr <= expr
        | expr = expr
        | not expr
        | (expr)
        | ID
        | integer
        | true
        | false

```

Figure 1: Syntax for the subset of *Cool* to be implemented in PA4

Essentially, you are leaving out language features that require implementing classes, including methods and method calls, attributes, inheritance, constructors, **new** and **case**. The class **Main** is the only class, and it is assumed to have a single method **Main.main() : Int**. This method becomes a simple, global LLVM function (We have give you special-purpose code to translate class *Main* and method *main()*). You only need to translate its body).

Note that the only types you must support are **Int** and **Bool**. **String** is not included because it requires objects, and **SELF.TYPE** is not needed in the absence of objects. To eliminate objects from the language, we need to make two small changes to the Cool typing rules:

1. You can assume that both branches of a conditional expression have the same type (both **Int** or both **Bool**). Therefore, the type of the whole expression will be either **Int** or **Bool**, and it will be the same as the type of the branches. In PA5, you will have to handle the case of *different* types in the branches, that are merged into the “join” (least upper bound) of the two types (see Section 7.5 of the *Cool* manual).

2. A loop expression has type `Int` and evaluates to the value 0. In PA5, you must implement the rule that a loop expression evaluates to a `void` value of type `Object` (Section 7.6 of the *Cool* manual). However, for PA4, `Object` is not supported.

The *Cool* runtime library is not used for PA4. Also, you cannot use class IO to perform input or output. Instead, you can return a result from the function `main`.

The code generator makes use of the AST constructed in PA2 and static analysis performed by the *semant* implemented in PA3. Your code generator should produce LLVM assembly code that faithfully implements any correct *Cool* program. There is no error recovery in code generation – all erroneous *Cool* programs have been detected by the front-end phases of the compiler.

This assignment gives you some flexibility in how exactly you generate LLVM code for individual *Cool* constructs. We will specify certain design choices in order to simplify the project (e.g., how to organize a virtual method table; how to interface to the the external system for printing and I/O). **You are responsible for other key design choices**, including how to organize the objects of each class, how to perform dynamic dispatch, and how to implement the basic built-in classes (`Int`, `Bool`, `String`, `IO`). Nevertheless, note that there are many key design goals to meet, and there are standard design approaches compilers use to meet these goals. We discuss these approaches in class or in this handout.

These two assignments will be considerably more difficult and take more time than the previous ones, so we suggest you get started right away. We have given you a simple exercise to help you learn the LLVM code representation and tools, before writing the code generator. The LLVM Programmer's Manual will also be useful for this purpose.

2 Files and Directories

To get started, download and unpack the *pa4* directory under Resources from the course website. This directory contains all the files that you will need for this PA. The subdirectory *pa4/src* contains the skeleton files for the code generation phase, which you will need to modify. You should not need to change any files in any other directories. The files you will need to modify are:

- `cgen.cc`

This file will contain your code generator. We have provided an implementation of some aspects of code generation; studying this code will help you write the rest of the code generator. It includes a call to code that will build an inheritance graph from the provided AST.

- `cgen.h`

This file is the header for the code generator. You may add anything you like to this file. It also provides classes for implementing the inheritance graph.

- `cool-tree.handcode.h`

This file contains the declarations of classes for AST nodes. You can add field declarations to the classes in *cool-tree.h* by editing *cool-tree.handcode.h*. The macros defined in there are included into *cool-tree.h* during definition. If you want to add a function `void foo(int)` as pure virtual into for example the *Feature_class* and its implementation into all subclasses, you have to put a definition `virtual void foo(int)=0` into *Feature_EXTRAS* and `void foo(int);` into *Feature_SHARED_EXTRAS*. The definitions of the methods should be added to *cgen.cc*.

- `value_printer.{h,cc}`, `operand.{h,cc}`

These files contain a small library for printing out simple LLVM instructions, which you may use for the assignment. The provided *cgen.h* includes *value_printer.h*. You may also print your assembly directly. Using the full LLVM API to assemble and pretty-print IR would require some changes to the makefile – let us know if you want to try this.

- `coolrt.{h,cc}`

These files provide an implementation for the *Cool* runtime library. You will not be using them for PA4.

- **README**

You should edit this file if you want your intermediate code generator to be scored with a different semantic analyzer than the submitted version. We will be using the reference lexer and parser for grading PA4 and PA5. If you want to resubmit your semantic analyzer (`semant.h`, `semant.cc`, `cool-tree.h`) and want us to use this version for grading your code generator, you should include updated files in `src` directory and put “2” in the first line of README. The default is “1” for using your submitted and scored semantic analyzer. We will use the reference semantic analyzer if you put “3” in the first line of README (10% automatic deduction from your score in this case).

To build the code generator, type

make cgen-1

in the directory `pa3/src`. This will link more files of support code to your directory and compile your files.

Note that `cgen.h` and `cgen.cc` use conditional compilation techniques (`ifdef` and `ifndef`) to build two different programs, depending on whether the symbol `PA4` is defined. We have set up the Makefile so that when you build `cgen-1`, `PA4` will not be defined, and these regions will not be compiled, nor will they appear in your binary `cgen-1`. You should not add any code to these sections for PA4 (you will need to later for PA5).

You should also take a short look at the other files in the `cool-support/src`, and `cool-support/include` directories. Especially `cgen_supp.cc` contains general support code for the code generator. You will find a number of handy functions here.

3 Designing the Code Generator

There are many possible ways to write the code generator. One reasonable strategy is to perform code generation in two passes; this is the strategy used by our solution and by the skeleton code. The first pass decides the object layout for each class, i.e. which LLVM data types to create for each class, and generates LLVM constants for all constants appearing in the program. Using this information, the second pass recursively walks each feature and generates LLVM code for each expression.

There are a number of things you must keep in mind while designing your code generator:

- You should have a clear picture of the runtime semantics of Cool programs. The semantics are described informally in the first part of the *CoolAid*, and a precise description of how Cool programs should behave is given in Section 12 of the manual.
- You should have a clear picture of LLVM instructions, types, and declarations.
- Think carefully about how and where objects, let-variables, and temporaries (intermediate values of expressions) are allocated in memory. The next section discusses this issue in some detail.
- You should generate unoptimized LLVM code, using a simple tree-walk similar to the one we discussed in class. Focus on generating reasonably efficient local code for each tree node, e.g., wherever possible, avoid extra casts, use `getelementptr` to index into objects (i.e., to compute the address of a structure field), use appropriate aggregate types, etc.
- Ignore the garbage collection requirement of Cool. You don't have to implement it. Just insert calls to the function `i8* @malloc(i32)` to allocate heap objects whenever needed, and never free these objects.

4 Representing Objects and Values in Cool

A major part of your compiler design is to develop the correct representation and memory allocation policies for objects and values in *Cool*, including explicit variables, heap objects, and temporaries. In PA4, you only need to be concerned with `Int` and `Bool` values, and only as variables or temporaries (not heap objects).

Here are the guidelines you should follow:

- Values of primitive types should be represented directly as virtual registers (of types `i32` and `i1`) in your generated code, always for PA4 and in most cases for PA5.
- The only time an `Int` or `Bool` must live on the heap in your compiler is if an actual object operation needs to be performed on it. Ordinary arithmetic operations (+, -, etc.) are not object operations. Assignment of a value to an `Int` or `Bool` variable is not an object operation. There are no object operations in PA4, and few in PA5. (Which are the object operations applicable to `Int` and `Bool`?)
- When an `Int` or `Bool` value must live on the heap (for an object operation), you must allocate a heap object for it and store the register value into that heap object before the object operation. This is called “boxing” the value.
- When an `Int` or `Bool` object value on the heap then needs to be used as a primitive type again (e.g., in an arithmetic operation), you must load the value out of the object into an LLVM register. This is called “unboxing” the value.
- One consequence of this strategy is that PA4 never needs an `Int` or `Bool` to be on the heap: it can always live in an LLVM register.
- A corollary is that the return value from `@Main.main` must be an `i32` and not `i32*`.
- Think of *let*-variables as names (pointers) for values (objects): this is the correct interpretation for *Cool* because the same variable can be assigned different values (and so must point to different heap objects) at different places within its *let*-block. Since a *let*-variable has a local scope, we can allocate it in the current stack frame using the `alloca` instruction. Even if a *let*-variable is of a primitive type (`i32` or `i1`), we will just allocate it on the stack and let *mem2reg* promote it to an SSA register for us. To enable this promotion, all `alloca` instructions should be placed in the entry of the *let*-block.

5 Testing the Code Generator

The directory *pa4/test-1* provides a place for you to test your code generator for PA4. *You should write your own test cases to test your compiler.* Use separate simple test initially, e.g., a single constant and simple arithmetic with two constants, and then work your way up to more complex expressions. A few days before the due date, we will provide a subset of our own tests.

The *test-1* directory contains its own Makefile. Some of the targets it provides are:

- `make file.ll`: compile the *Cool* program *file.cl* to LLVM assembly
- `make file.bc`: create an LLVM bytecode file from *file.ll*
- `make file.exe`: create a linked executable from *file.bc* (this is no-op for PA4)
- `make file.out`: execute *file.exe* and put the output in *file.out*
- `make file.verify`: verify your LLVM code obeys LLVM language rules
- `make file-opt.bc`: create an optimized LLVM bytecode file from *file.exe.bc*. This is just so you can see whether your code can be optimized effectively by available techniques in LLVM.

Note that Makefile here uses the reference semantic analyzer by default for compilation. To test your version of semantic analyzer along with the code generator, you should modify *Makefile.common* under *pa4* directory to define `SEMANT` as a path to your semantic analyzer.

To be sure that you generate correct LLVM code you should call the LLVM verification path with every program that you generate. You can do this by saying `make file.verify` as described above. See the target `%.verify` in *pa4/Makefile.common* for the command used.

In order to run a *Cool* program and inspect its result, your compiler should add a `main()` function to the generated LLVM module. This function should call `@Main.main()` and print the result. It should look like the following or equivalent:

```
define i32 @main() {
entry:
    %tpm.0 = call i32 @Main_main( )
    %tpm.1 = getelementptr [25 x i8], [25 x i8]* @.str, i32 0, i32 0
    %tpm.2 = call i32(i8*, ... ) @printf( i8* %tpm.1, i32 %tpm.0 )
    ret i32 0
}
```

The `.exe` target will fail until your compiler generates a valid assembly file that defines a `main` function with the right signature, and `.out` targets may fail rather than capture output if your generated assembly has sufficiently serious errors.

You should generate this function explicitly using LLVM IR features. To make this easier for you, we've provided a skeleton routine called `CgenClassTable::code main()` in `cgen.cc`.

Your code generation phase executable `cgen` takes a `-c` flag to generate debugging information. This is set whenever you define `debug` true in your Makefile (the default). Using this flag merely causes `cgen_debug` (a global variable) to be set. Adding the actual code to produce useful debugging information is up to you. See the project README for details.

6 Notes

Since writing a code generator is a fairly big task, we suggest that you take the following steps in order to build your compiler. Make sure to test each portion of code as you complete it.

1. Start by generating the function `@main()` as described above, so that you can test your compiler even in the early stages of your work.
2. Start by implementing `Int` and `Bool` constants. At first, you can just generate them as the `i32` and `i1` LLVM primitives. *Test your compiler!*
3. Once you have constants, you can implement arithmetic and comparison operators. You can also implement block expressions at this time (e.g., `1 + 2`; `2 <= 1`). *Test your compiler!*
4. Now try implementing `let`, following the allocation guidelines in the previous section. Use the environment to keep track of the binding from *Cool* variable names to memory locations (i.e., to LLVM `alloca/global/malloc` values). *You know what you need to do now!*
5. Next, implement assignment. Here, you will need to think about how the LHS and RHS are implemented, and what should be copied over.
6. Next, tackle `loop` and `if-then-else`. For these, you will need to learn more about LLVM `BasicBlocks`. In addition to the regular sources, the *Stacker* documentation contains some useful tips on using `BasicBlocks` to implement control structures.

The result of an `if-then-else` is a merge of the results of the two branches. You can allocate an `i32` or `i1` in the stack (depending on the type of the then-else branches) and then store a different result in each of the branches.

7. The final step: implement runtime error handling, if you haven't already. There are only a few cases you need to check, and they're listed in the back of the *Cool* manual. For PA4, the only possible error is divide-by-zero. Your program should call the function `abort()` if this happens. We have given you code to insert a declaration for `abort()` in the module.
8. Now test your compiler more thoroughly. You can use the *Cool* files we will give you in the `test-1` directory, but you should also make your own tests to stress individual cases.

7 What and How to Turn In

You have to turn in the *pa4* directory with your modified version of *coo-tree.handcode.h*, *cgen.cc*, and *cgen.h* (and *README* optionally) after compressing it with

```
tar -cvf pa4-[your student id].tar pa4
```

Make sure all your code for the semantic analyzer is in *cool-tree.h* (and/or *coo-tree.handcode.h*), *semant.h*, and *semant.cc*. You can upload the compressed file to the board for assignment submission at the course website. Please do not copy or modify any part of the support code. The provided files are the ones that will be used in the grading process.